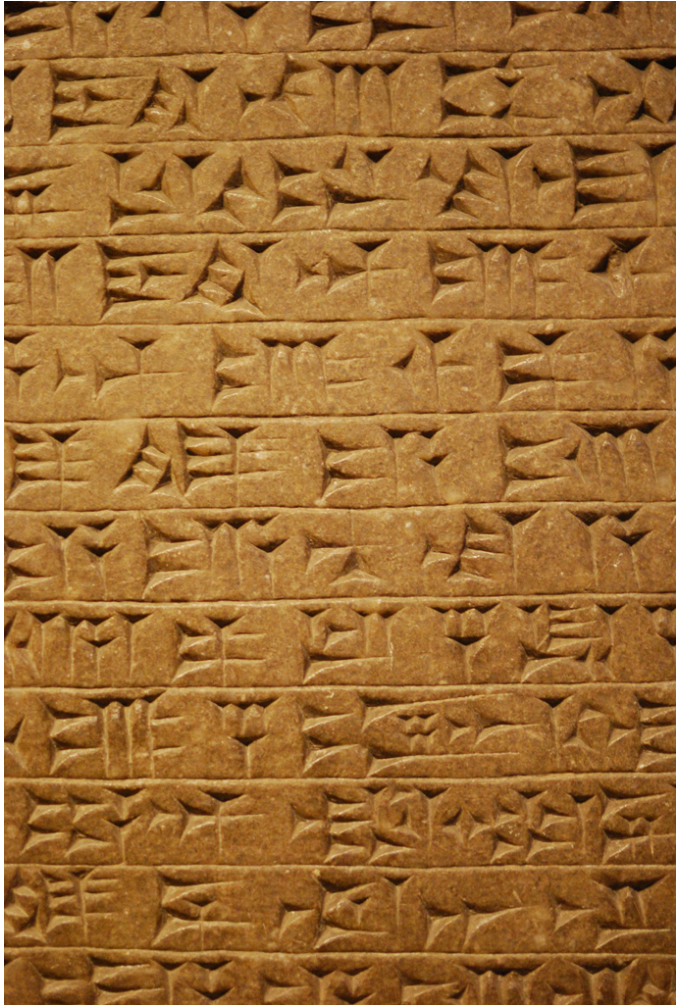


History

This is how one pictures the angel of history. His face is turned toward the past. Where we perceive a chain of events, he sees one single catastrophe which keeps piling wreckage and hurls it in front of his feet. The angel would like to stay, awaken the dead, and make whole what has been smashed. But a storm is blowing in from Paradise; it has got caught in his wings with such a violence that the angel can no longer close them. The storm irresistibly propels him into the future to which his back is turned, while the pile of debris before him grows skyward. This storm is what we call progress.

(Walter Benjamin, 1940)

The Dawn of Computation



Babylonian
cuneiform
circa 2000 B.C.
(Photo by Matt Neale)

Early Computation


“computer” = human being performing computation

Euclid’s algorithm circa 300 B.C.

Abu Ja’far Muhammad ibn Musa Al-Khwarizmi’s books on algebra and arithmetic computation using Indo-Arabic numerals, circa 800 A.D.

Isaac Newton (1643-1727)

Charles Babbage (1791-1871)



Difference Engine (1819)

Analytical Engine (1834)

Mechanical computation for
military applications

The specification of
computational operations
was separated from their
execution

Babbage's designs were
technically too ambitious

Ada Augusta Byron (1815-1852)



Assisted Babbage in explaining and promoting his ideas

Wrote articles describing the operation and use of the Analytical Engine

The first computer scientist?

David Hilbert (1862-1943)

Formalized the axiomatic treatment of
Euclidean geometry

Hilbert's 23 problems (ICM, 1900)

Problem #2: Is mathematics consistent?

The meaning of proof

Axiom: $\forall n : n + 0 = n.$

Math statement: “The square of any even number is even.”

Formula: $\forall n (\exists k : n = k + k \Rightarrow \exists m : m + m = n * n)$

Proof: Finite sequence of axioms (basic true statements) and derivations of new true statements (e.g. ϕ and $\phi \rightarrow \sigma$ yield σ).

Theorem: A mathematical statement ϕ together with a proof deriving ϕ within a given system of axioms and derivation rules.

Hilbert's questions (1920's)

Is mathematics complete? Meaning: for any formula ϕ , if ϕ is true, then ϕ is provable.

Is mathematics consistent? Meaning: for any formula ϕ , there aren't proofs of both ϕ and $\neg\phi$.

Is there a procedure to, given a formula ϕ , produce a proof of ϕ , or show there isn't one?

Hilbert believed the answers would be “yes”.

Kurt Gödel (1906-78)

Gödel's answers to Hilbert (1929-30):

Any axiom system powerful enough to describe arithmetic on integers is not complete.

If it is consistent, its consistency cannot be proved within the system.

Sketch of Gödel's proof

Define a mapping between logical formulas and numbers.

Use it to define mathematical statements saying “This number represents a valid formula”, “This number represents a sequence of valid formulae”, “This number represents a valid proof”, “This number represents a provable formula”.

Construct a formula ϕ represented by a number n that says “The formula represented by n is not provable”. The formula ϕ cannot be false, so it must be true but not provable.

What remained of Hilbert's questions

Is there a procedure which, given a formula ϕ , either proves ϕ , shows it false, or correctly concludes ϕ is not provable?

The answer to this requires a precise definition of “a procedure”, in other words, a formal model of computation.

Alonzo Church (1903-1995)

Set out to give a final “no” answer to this last question

With his student Kleene, created notation to describe functions on the natural numbers.

Church and Kleene's notation

They wanted to modify Russell and Whitehead's notation for the class of all x satisfying a predicate f : $\hat{x} f(x)$.

But their notion was somewhat different, so they tried putting the caret before: $\hat{} x$.

Their typewriter could not type this, but had Greek letters.

Perhaps a capital lambda? Λx .

Too much like the symbol for logical AND: \wedge .

Perhaps a lower-case lambda? λx .

The lambda calculus

The function that added 2 to its argument would be represented by $\lambda x.x + 2$.

The function that subtracted its second argument from its first would be written $\lambda x.\lambda y.x - y$.

fx applies function f to argument x .

fxy means $(fx)y$ (left-associativity).

To prove something is impossible to express in some notation, the notation should be as simple as possible.

To make things even simpler, the lambda calculus did not permit naming of functions (only parameters), naming of constants like 2, or naming of functions like +.

It had three grammar rules and one reduction rule (function application).

How could it say anything at all?

Numbers from nothing (Cantor-style)

$0 \equiv \emptyset$ or $\{\}$ (the empty set)

$1 \equiv \{\emptyset\}$

$2 \equiv \{\{\emptyset\}, \emptyset\}$

In general, n is represented by the set containing the sets representing $n - 1, n - 2, \dots, 0$.

This is the way that arithmetic can be built up from the basic axioms of set theory.

Numbers from nothing (Church-style)

$0 \equiv \lambda f. \lambda x. x$ (the function which ignores its argument and returns the identity function)

$1 \equiv \lambda f. \lambda x. f x$ (the function which, when given as argument a function f , returns the same function).

$2 \equiv \lambda f. \lambda x. f(f x)$ (the function which, when given as argument a function f , returns f composed with itself or $f \circ f$).

In general, n is the function which does n -fold composition.

With some care, one can write down short expressions for the addition and multiplication functions.

Similar ideas will create Boolean values, logical functions, and conditional expressions.

General recursion without naming is harder, but still possible.

The lambda calculus is a general model of computation.

Church's proof

Church proved that there was no computational procedure to tell if two lambda expressions were equivalent (represented the same function).

His proof mirrored Gödel's, using a way of encoding lambda expressions using numbers, and provided a “no” answer to the idea of deciding provability of formulae.

This was published in 1936.

Independently, a few months later, a British mathematician came up with a simpler proof.


Alan Turing (1912-1954)

Turing defined a different model of computation, and chose a different problem to prove uncomputable.

This resulted in a simpler and more influential proof.

Turing's model of computation

f: state char \rightarrow state char move



Finite state control plus unbounded storage tape

Turing's proof (1936)

Turing showed how to code the finite state control using characters.

He then assumed that there was a machine that could process such a description and tell whether the coded machine would halt or not when fed its own description as an input.

Using this machine, one can define a second machine that acts on this information.

The second machine uses the first machine to see if its input represents a coded machine which halts when fed its own description.

If so, the second machine runs forever; otherwise, it halts.

Feeding the description of the second machine to itself creates a contradiction: it halts iff it doesn't halt.

So the first machine cannot exist.

Turing's proof also demonstrates the undecidability of proving formulae.

Advantages of Turing's ideas

Turing's ideas can be adapted to give a similar proof in the lambda calculus model.

Upon learning of Church's work, Turing quickly sketched the equivalence of the two models.

Turing's model bears a closer resemblance to an intuitive idea of real computation.

It would influence the future development of hardware and thus software, even though reasoning about programs is more difficult in it.

Turing went to America to study with Church at Princeton, earning his PhD in 1939.

During the war, he was instrumental in an effort to break encrypted German radio traffic, resulting in the development of what we now know to be the world's first working electronic computer (Colossus).

Turing made further contributions to hardware and software design in the UK, and to the field of artificial intelligence, before his untimely death in 1954.

John von Neumann (1903-1957)

von Neumann was a founding member of the Institute for Advanced Study at Princeton.

In 1946 he visited the developers of ENIAC at the University of Pennsylvania, and wrote an influential “Report on the EDVAC” regarding its successor.

Features: random-access memory, CPU, fetch-execute loop, stored program.

Lacking: support for recursion (unlike Turing’s UK designs)

Grace Murray Hopper (1906-1992)

Wrote first compiler, defined first English-like data processing language

Ideas later folded into COBOL (1959)



FORTRAN (1957)

Early programming language influenced by architecture.

```
      INTEGER FN, FNM1, TEMP
      FN = 1
      FNM1 = 0
      DO 20 I = 1, 10, 1
      PRINT 10, I, FN
10    FORMAT(I3, 1X, I3)
      TEMP = FN + FNM1
      FNM1 = FN
20    FN = TEMP
```

FORTRAN was designed by John Backus, and became the dominant language for numerical and scientific computation. Backus also invented a notation for language description that is popular in programming language design today.

Backus won the Turing Award in 1978, and used the associated lecture to criticize the continued dominance of von Neumann's architectural model and the programming languages inspired by it.

He proposed a functional programming language for parallel/distributed computation.

FORTRAN and COBOL, reflecting the Turing - von Neumann approach, dominated practical computing through most of the '60's and '70's.

Many other computer languages were defined, enjoyed brief and modest success, and then were forgotten.

Church's work proved useful in the field of operational semantics, which sought to treat the meaning of programs mathematically.

It also was inspirational in the design of a still-popular high-level programming language called Lisp.

John McCarthy (1927-2011)

McCarthy, an AI researcher at MIT, was frustrated by the inexpressiveness of machine languages and the primitive programming languages arising from them (no recursion, no conditional expressions).

In 1958, he designed and implemented Lisp (LISt Processor), taking ideas from the lambda calculus and the theory of recursive functions.

His 1960 paper on Lisp described the core of the language in terms that CS 135 students would recognize.

McCarthy's Lisp

McCarthy defined these primitive functions: `atom` (the negation of `cons?`), `eq`, `car` (`first`), `cdr` (`rest`), and `cons`.

He also defined the special forms `quote`, `lambda`, `cond`, and `label` (`define`).

Using these, he showed how to build many other useful functions.

The evolution of Lisp

The first implementation of Lisp, on the IBM 704, could fit two machine addresses (15 bits) into parts of one machine word (36 bits) called the address and decrement parts. Machine instructions facilitated such manipulation.

This led to the language terms `car` and `cdr` which persist in Racket and Lisp to this day.

Lisp quickly evolved to include proper numbers, input/output, and a more comprehensive set of built-in functions.

Lisp became the dominant language for artificial intelligence implementations.

It encouraged redefinition and customization of the language environments, leading to a proliferation of implementations.

It also challenged memory capabilities of 1970's computers, and some special-purpose "Lisp machines" were built.

Modern hardware is up to the task, and the major Lisp groups met and agreed on the Common Lisp standard in the 1980's.

The origins of Scheme

In 1976, researcher Carl Hewitt designed an ambitious Lisp-like language called Planner, with facilities for automated backtracking for goal search.

Researcher Gerald Sussman, with others, created Conniver, a version which made the control structures more explicit and more efficient.

Hewitt also proposed the *actors* theory of distributed object-oriented computation, with facilities to move computation from one “actor” to another.

Sussman and his graduate student Guy Steele were having trouble understanding some of the implications of the actors model, so they wrote a small Lisp interpreter for a “toy” actors language.

Sussman suggested using lexical scoping instead of the dynamic scoping used by Lisp, and a single namespace (Lisp has separate namespaces for functions and variables).

The migration of computation was captured by the idea of *continuations*, which made the unevaluated part of a computation into a manipulable object.

Sussman and Steele realized that function creation and actor creation were almost identical, as were function application, message passing, and actor invocation.

This simplified their language even further.

They wanted to call it Schemer, but there was a six-character file name limit on their system, so it became Scheme.

Research groups at other universities started using Scheme to study programming languages.

Sussman and Steele today

Steele went to work for Thinking Machines and then Sun Microsystems, where he became an early member of the team developing the Java programming language.

Sussman, together with colleague Hal Abelson, started using Scheme in the undergraduate program at MIT. Their textbook, “Structure and Interpretation of Computer Programs”, is used there, as well as at many other major institutions.

The authors of the HtDP textbook developed an extension of Scheme (PLT Scheme) and its learning environment (DrScheme) to remedy the following perceived deficiencies of SICP:

- lack of programming methodology
- complex domain knowledge required
- steep, frustrating learning curve
- insufficient preparation for future courses

As PLT Scheme and the teaching languages diverged further from Sussman and Steele's Scheme, they renamed their language Racket in 2010.

Summing up CS 135

With only a few language constructs (**define**, **cond**, **define-struct**, **cons**, **local**, **lambda**) we have described and implemented ideas from introductory computer science in a brief and natural manner.

We have done so without many of the features (static types, mutation, I/O) that courses using conventional languages have to introduce on the first day. The ideas we have covered carry over into languages in more widespread use.

We hope you have been convinced that the goal of computer science is to implement useful computation in a way that is correct and efficient as far as the machine is concerned, but that is understandable and extendable as far as other humans are concerned.

These themes will continue in CS 136, but new themes will be added, and a new programming language using a different paradigm (though we will continue to use Racket as well).

Looking ahead to CS 136

We have been fortunate to work with very small languages (the teaching languages) writing very small programs which operate on small amounts of data.

In CS 136, we will broaden our scope, moving towards the messy but also rewarding realm of the “real world”.

The main theme of CS 136 is scalability: what are the issues which arise when things get bigger, and how do we deal with them?

How do we organize a program that is bigger than a few screenfuls?

How do we reuse and share code, apart from cutting-and-pasting it into a new program file?

How do we design programs so that they run efficiently?

What changes might be necessary to our notion of types and to the way we handle errors when there is a much greater distance in time and space between when the program is written and when it is run?

When is it appropriate to abstract away from implementation details for the sake of the big picture, and when must we focus on exactly what is happening at lower levels for the sake of efficiency?

These are issues which arise not just for computer scientists, but for anyone making use of computation in a working environment.

We can build on what we have learned this term in order to meet these challenges with confidence.